

# Machine Learning – Final Increment (Project)

CS 5710 (CRN 22002)

Project Title: A Comparison of Machine Learning Algorithms on Detection of Phishing Websites

Team Members:

1. Kamala Ramesh – 700745451
2. Madhuri Vedula – 700741625
3. Priyanka Bojja – 700739528
4. Sai Priyanka Narra – 700741613

Please find below the screenshots of the implementation

```
In [1]: import pandas as pd
import itertools
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import xgboost as xgb
import os
import seaborn as sns
from wordcloud import WordCloud
```

```
In [2]: from google.colab import drive
drive.mount('/content/gdrive')

Mounted at /content/gdrive
```

```
In [3]: !pip install wordcloud
```

```
In [4]: df=pd.read_csv('/content/gdrive/My Drive/data_updated_phish.csv')
print(df.shape)
df.head()
```

(165198, 2)

```
Out[4]:
```

	url	type
0	br-icloud.com.br	phishing
1	signin.eby.de.zukrueyxcztzmmqi.civpro.co.za	phishing
2	http://www.marketingbyinternet.com/mo/e56508df...	phishing
3	https://docs.google.com/spreadsheet/viewform?f...	phishing
4	retajconsultancy.com	phishing

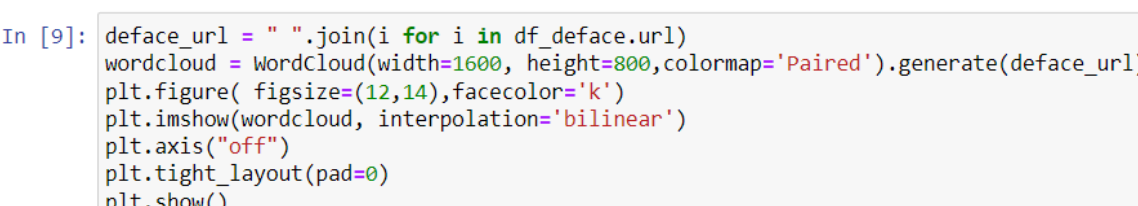
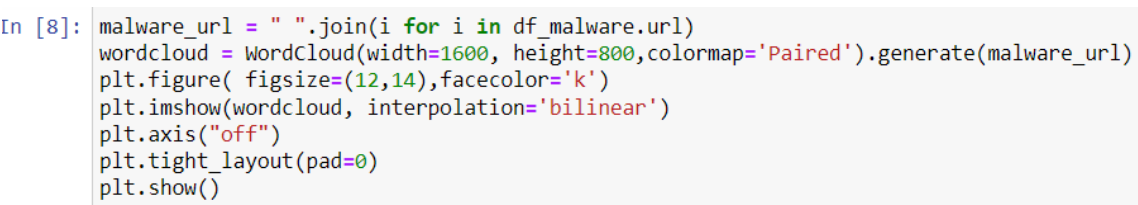
```
In [5]: df.type.value_counts()
```

```
Out[5]: benign      101612
defacement    23709
phishing      22942
malware       16935
Name: type, dtype: int64
```

## Word Cloud

```
In [6]: df_phish = df[df.type=='phishing']
df_malware = df[df.type=='malware']
df_deface = df[df.type=='defacement']
df_benign = df[df.type=='benign']
```

```
In [7]: phish_url = " ".join(i for i in df_phish.url)
wordcloud = WordCloud(width=1600, height=800, colormap='Paired').generate(phish_url)
plt.figure( figsize=(12,14),facecolor='k')
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.tight_layout(pad=0)
plt.show()
```







## Feature Engineering

```
In [11]: import re
#Use of IP or not in domain
def having_ip_address(url):
    match = re.search(
        '([01]?\\d\\d?[2[0-4]\\d|25[0-5])\\.([01]?\\d\\d?[2[0-4]\\d|25[0-5])\\.([01]?\\d\\d?[2[0-4]\\d|25[0-5])\\.([01]?\\d\\d?[2[0-4]\\d|25[0-5])\\/'
        '([01]?\\d\\d?[2[0-4]\\d|25[0-5])\\/' # IPv4
        '([0x[0-9a-fA-F]{1,2})\\.([0x[0-9a-fA-F]{1,2})\\.([0x[0-9a-fA-F]{1,2})\\.([0x[0-9a-fA-F]{1,2})\\/' # IPv4 in hexadecimal
        '([a-fA-F0-9]{1,4}){7}[a-fA-F0-9]{1,4}', url) # Ipv6
    if match:
        # print match.group()
        return 1
    else:
        # print 'No matching pattern found'
        return 0
df['use_of_ip'] = df['url'].apply(lambda i: having_ip_address(i))
```

```
In [12]: from urllib.parse import urlparse

def abnormal_url(url):
    hostname = urlparse(url).hostname
    hostname = str(hostname)
    match = re.search(hostname, url)
    if match:
        # print match.group()
        return 1
    else:
        # print 'No matching pattern found'
        return 0

df['abnormal_url'] = df['url'].apply(lambda i: abnormal_url(i))
```

```
In [13]: #!pip install googlesearch-python
```

```
In [14]: from googlesearch import search
```

```
In [15]: def google_index(url):
    site = search(url, 5)
    return 1 if site else 0
df['google_index'] = df['url'].apply(lambda i: google_index(i))
```

```
In [16]: def count_dot(url):
    count_dot = url.count('.')
    return count_dot

df['count.'] = df['url'].apply(lambda i: count_dot(i))
df.head()
```

```
Out[16]:
```

	url	type	use_of_ip	abnormal_url	google_index	count.
0	br-icloud.com.br	phishing	0	0	1	2
1	signin.eby.de.zukruxgxtzmmqi.civpro.co.za	phishing	0	0	1	6
2	http://www.marketingbyinternet.com/mo/e56508df...	phishing	0	1	1	2
3	https://docs.google.com/spreadsheet/viewform?f...	phishing	0	1	1	2
4	retajconsultancy.com	phishing	0	0	1	1

```

In [17]: def count_www(url):
          url.count('www')
          return url.count('www')

df['count-www'] = df['url'].apply(lambda i: count_www(i))

def count_atrate(url):
    return url.count('@')

df['count@'] = df['url'].apply(lambda i: count_atrate(i))

def no_of_dir(url):
    urldir = urlparse(url).path
    return urldir.count('/')

df['count_dir'] = df['url'].apply(lambda i: no_of_dir(i))

def no_of_embed(url):
    urldir = urlparse(url).path
    return urldir.count('///')

df['count_embed_domian'] = df['url'].apply(lambda i: no_of_embed(i))

def shortening_service(url):
    match = re.search('bit\.ly|goo\.gl|shorte\.st|go2l\.ink|x\.co|ow\.ly|t\.co|tinyurl|tr\.im|is\.gd|cli\.gs|'
                      'yfrog\.com|migre\.me|ff\.im|tiny\.cc|url4\.eu|twit\.ac|su\.pr|twurl\.nl|snipurl\.com|'
                      'short\.to|BudURL\.com|ping\.fm|post\.ly|Just\.as|bkite\.com|snipr\.com|fic\.kr|loopt\.us|'
                      'doiop\.com|short\.ie|kl\.am|wp\.me|rubyurl\.com|om\.ly|to\.ly|bit\.do|t\.co|lnkd\.in|'
                      'db\.tt|qr\.ae|adf\.ly|goo\.gl|bitly\.com|cur\.lv|tinyurl\.com|ow\.ly|bit\.ly|ity\.im|'
                      'q\.gs|is\.gd|po\.st|bc\.vc|twitthis\.com|u\.to|j\.mp|buzurl\.com|cutt\.us|u\.bb|yourls\.org|'
                      'x\.co|prettylinkpro\.com|scrnch\.me|filoops\.info|vzturl\.com|qr\.net|1url\.com|tweez\.me|v\.gd|'
                      'tr\.im|link\.zip\.net',
                      url)

```

```
    if match:
        return 1
    else:
        return 0
```

```
df['short_url'] = df['url'].apply(lambda i: shortening_service(i))
```

```
In [18]: def count_https(url):
          return url.count('https')

df['count-https'] = df['url'].apply(lambda i : count_https(i))

def count_http(url):
    return url.count('http')

df['count-http'] = df['url'].apply(lambda i : count_http(i))
```

```
In [19]: def count_per(url):
          return url.count('%')

df['count%'] = df['url'].apply(lambda i : count_per(i))

def count_ques(url):
    return url.count('?')

df['count?'] = df['url'].apply(lambda i: count_ques(i))

def count_hyphen(url):
    return url.count('-')

df['count-'] = df['url'].apply(lambda i: count_hyphen(i))

def count_equal(url):
    return url.count('=')
```

```

df['count='] = df['url'].apply(lambda i: count_equal(i))
def url_length(url):
    return len(str(url))

#Length of URL
df['url_length'] = df['url'].apply(lambda i: url_length(i))
#Hostname Length

def hostname_length(url):
    return len(urlparse(url).netloc)

df['hostname_length'] = df['url'].apply(lambda i: hostname_length(i))

def suspicious_words(url):
    match = re.search('PayPal|login|signin|bank|account|update|free|lucky|service|bonus|ebayisapi|webscr',
        url)
    if match:
        return 1
    else:
        return 0
df['sus_url'] = df['url'].apply(lambda i: suspicious_words(i))

def digit_count(url):
    digits = 0
    for i in url:
        if i.isnumeric():
            digits = digits + 1
    return digits

df['count-digits']= df['url'].apply(lambda i: digit_count(i))

```

```

def letter_count(url):
    letters = 0
    for i in url:
        if i.isalpha():
            letters = letters + 1
    return letters

df['count-letters']= df['url'].apply(lambda i: letter_count(i))

df.head()

```

Out[19]:

	url	type	use_of_ip	abnormal_url	google_index	count.	count- www	count@	count_dir	count_embed_domian	...
0	br-icloud.com.br	phishing	0	0	1	2	0	0	0	0	...
1	signin.eby.de.zukruxgxtzmmqi.civpro.co.za	phishing	0	0	1	6	0	0	0	0	...
2	http://www.marketingbyinternet.com/mo/e56508df...	phishing	0	1	1	2	1	0	3	0	...
3	https://docs.google.com/spreadsheet/viewform?f...	phishing	0	1	1	2	0	0	2	0	...
4	retajconsultancy.com	phishing	0	0	1	1	0	0	0	0	...

5 rows x 22 columns





```
In [20]: !pip install tld

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting tld
  Downloading tld-0.13-py2.py3-none-any.whl (263 kB)
    263.8/263.8 kB 17.5 MB/s eta 0:00:00
Installing collected packages: tld
Successfully installed tld-0.13
```

```
In [21]: #Importing dependencies
from urllib.parse import urlparse
from tld import get_tld
import os.path

#First Directory Length
def fd_length(url):
    urlpath= urlparse(url).path
    try:
        return len(urlpath.split('/')[1])
    except:
        return 0

df['fd_length'] = df['url'].apply(lambda i: fd_length(i))

#Length of Top Level Domain
df['tld'] = df['url'].apply(lambda i: get_tld(i,fail_silently=True))

def tld_length(tld):
    try:
        return len(tld)
    except:
        return -1

df['tld_length'] = df['tld'].apply(lambda i: tld_length(i))
```

```
In [22]: df = df.drop("tld",1)

<ipython-input-22-beaedf415f7f>:1: FutureWarning: In a future version of pandas all arguments of DataFrame.drop except for the
argument 'labels' will be keyword-only.
df = df.drop("tld",1)
```


```
In [23]: df.columns
```

```
Out[23]: Index(['url', 'type', 'use_of_ip', 'abnormal_url', 'google_index', 'count.',
               'count-www', 'count@', 'count_dir', 'count_embed_doman', 'short_url',
               'count-https', 'count-http', 'count%', 'count?', 'count-', 'count=',
               'url_length', 'hostname_length', 'sus_url', 'count-digits',
               'count-letters', 'fd_length', 'tld_length'],
              dtype='object')
```

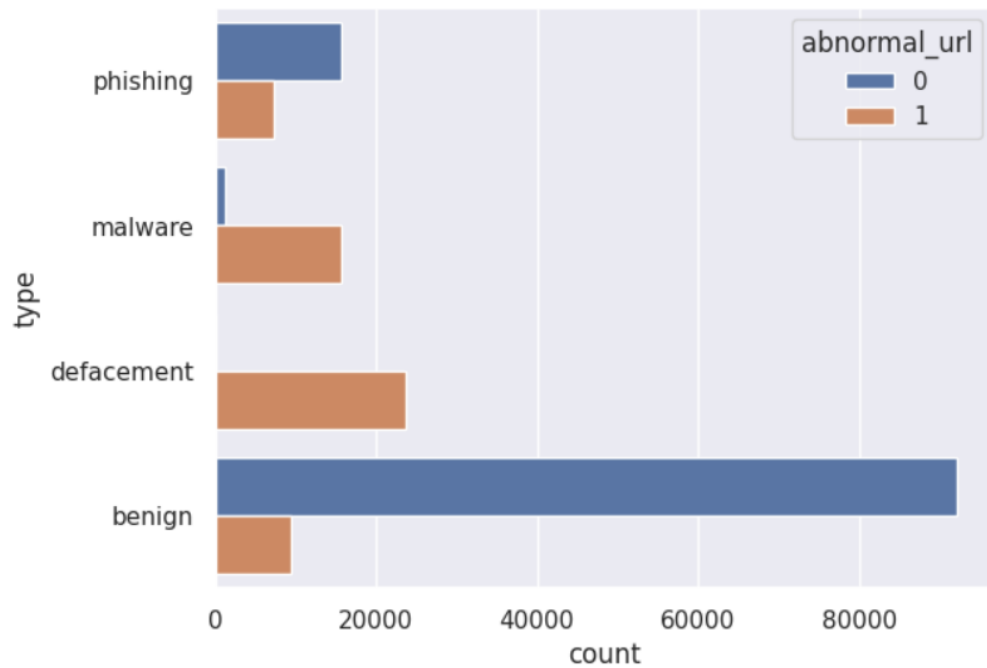
```
In [24]: df['type'].value_counts()
```

```
Out[24]: benign      101612
defacement    23709
phishing      22942
malware       16935
Name: type, dtype: int64
```


```
In [25]: import seaborn as sns
sns.set(style="darkgrid")
ax = sns.countplot(y="type", data=df,hue="use_of_ip")
```




```
In [26]: sns.set(style="darkgrid")
ax = sns.countplot(y="type", data=df, hue="abnormal_url")
```




```
In [27]: sns.set(style="darkgrid")
ax = sns.countplot(y="type", data=df, hue="google_index")
```




```
In [28]: sns.set(style="darkgrid")
ax = sns.countplot(y="type", data=df, hue="short_url")
```



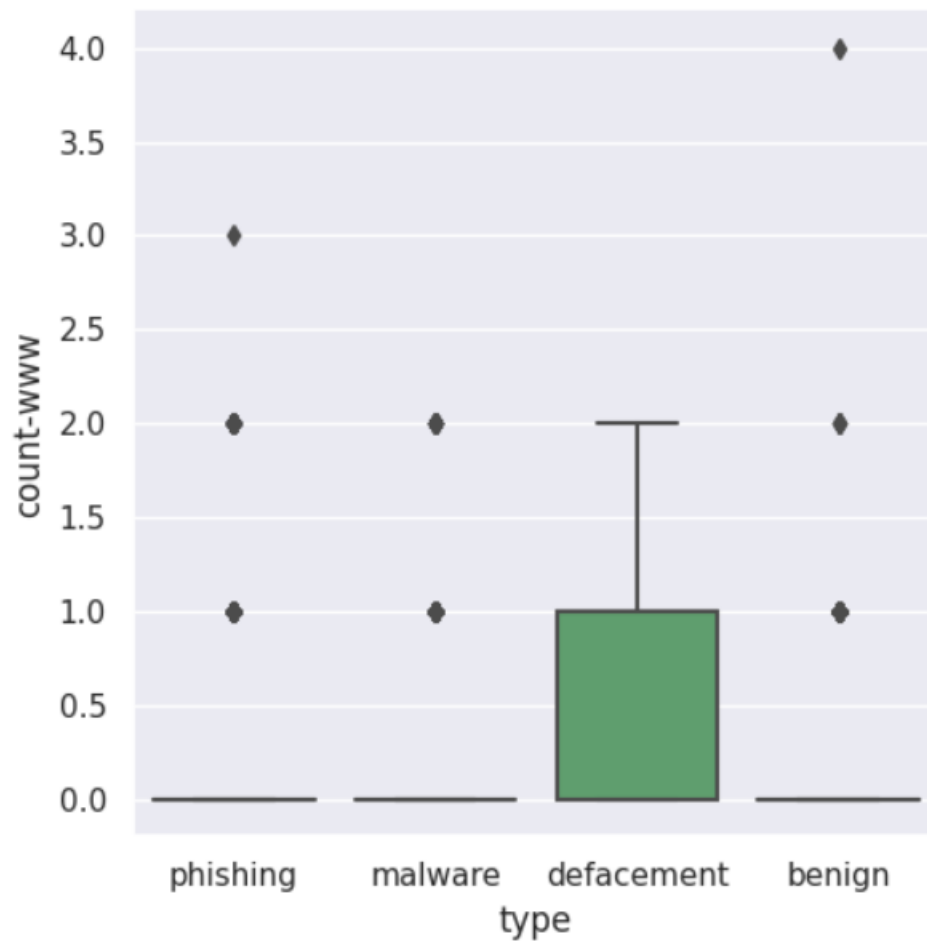
```
In [29]: sns.set(style="darkgrid")  
ax = sns.countplot(y="type", data=df, hue="sus_url")
```



```
In [30]: sns.set(style="darkgrid")
ax = sns.catplot(x="type", y="count.", kind="box", data=df)
```




```
In [31]: sns.set(style="darkgrid")
ax = sns.catplot(x="type", y="count-www", kind="box", data=df)
```

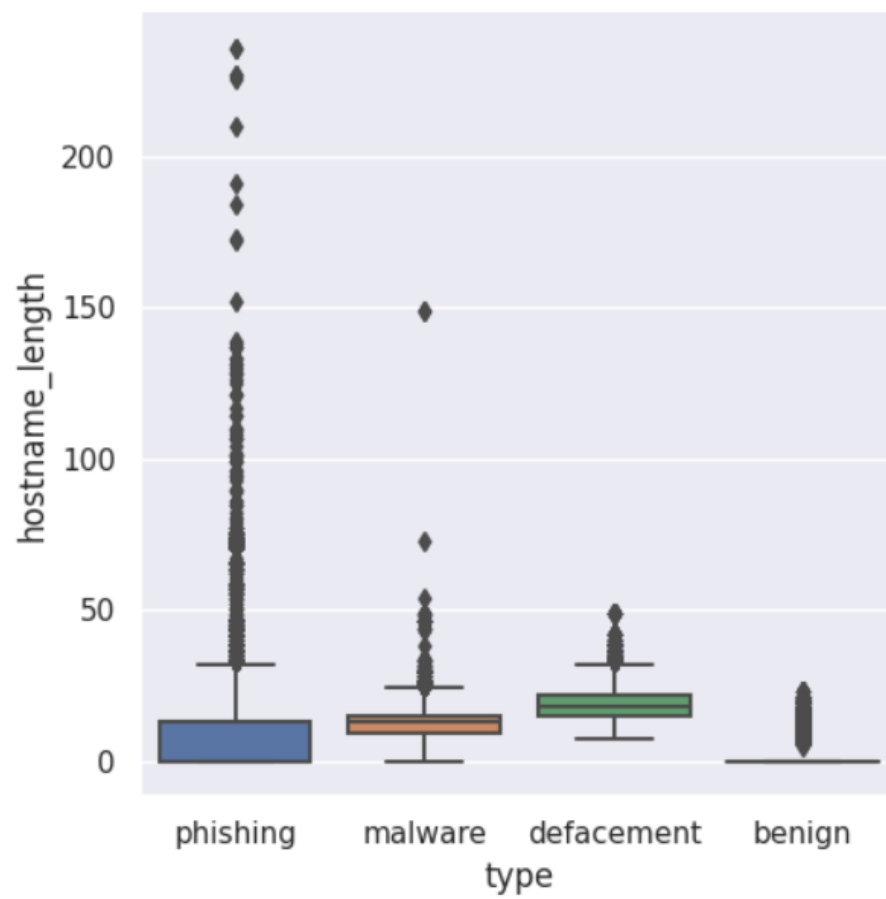





```
In [32]: sns.set(style="darkgrid")
ax = sns.catplot(x="type", y="count_dir", kind="box", data=df)
```




```
In [33]: sns.set(style="darkgrid")
ax = sns.catplot(x="type", y="hostname_length", kind="box", data=df)
```



```
In [34]: sns.set(style="darkgrid")
ax = sns.catplot(x="type", y="fd_length", kind="box", data=df)
```



```
In [35]: sns.set(style="darkgrid")
ax = sns.catplot(x="type", y="tld_length", kind="box", data=df)
```



## Target Encoding

```
In [36]: from sklearn.preprocessing import LabelEncoder

lb_make = LabelEncoder()
df["type_code"] = lb_make.fit_transform(df["type"])
df["type_code"].value_counts()
```

```
Out[36]: 0    101612
         1     23709
         3     22942
         2     16935
         Name: type_code, dtype: int64
```

## Test Train split

```
In [37]: #Predictor Variables
# filtering out google_index and use of IP as it has only 1 value
X = df[['use_of_ip', 'abnormal_url', 'count.', 'count-www', 'count@',
        'count_dir', 'count_embed_domian', 'short_url', 'count-https',
        'count-http', 'count%', 'count?', 'count-', 'count=', 'url_length',
        'hostname_length', 'sus_url', 'fd_length', 'tld_length', 'count-digits',
        'count-letters']]

#Target Variable
y = df['type_code']
```

```
In [38]: X.head()
```

```
Out[38]:
```

	use_of_ip	abnormal_url	count.	count-www	count@	count_dir	count_embed_domian	short_url	count-https	count-http	...	count?	count-	count=	url_length	hostname
0	0	0	0	2	0	0	0	0	0	0	...	0	1	0	16	
1	0	0	0	6	0	0	0	0	0	0	...	0	0	0	42	
2	0	1	2	1	0	3	0	1	0	1	...	0	0	0	71	
3	0	1	2	0	0	2	0	0	1	1	...	1	0	1	87	
4	0	0	1	0	0	0	0	0	0	0	...	0	0	0	20	

5 rows × 17 columns

```
In [39]: X.columns
```

```
Out[39]: Index(['use_of_ip', 'abnormal_url', 'count.', 'count-www', 'count@',
               'count_dir', 'count_embed_domian', 'short_url', 'count-https',
               'count-http', 'count%', 'count?', 'count-', 'count=', 'url_length',
               'hostname_length', 'sus_url', 'fd_length', 'tld_length', 'count-digits',
               'count-letters'],
              dtype='object')
```

```
In [40]: X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, shuffle=True, random_state=5)
```

```
In [41]: print(X_train.shape)
print(X_test.shape)
```

```
(132158, 21)
(33040, 21)
```

## Random Forest

```
In [42]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators=100, max_features='sqrt')
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print(classification_report(y_test, y_pred_rf, target_names=['benign', 'defacement', 'phishing', 'malware']))

score = accuracy_score(y_test, y_pred_rf)
print("accuracy: %.3f" % score)
```

	precision	recall	f1-score	support
benign	1.00	1.00	1.00	20323
defacement	0.97	0.99	0.98	4742
phishing	0.99	0.92	0.95	3387
malware	0.95	0.97	0.96	4588
accuracy			0.99	33040
macro avg	0.98	0.97	0.97	33040
weighted avg	0.99	0.99	0.99	33040


```
accuracy: 0.985
```

```
In [57]: # calculate precision, recall, and F1 score
from sklearn.metrics import precision_score, recall_score, f1_score
precision = precision_score(y_test, y_pred_rf, average="weighted")
recall = recall_score(y_test, y_pred_rf, average="weighted")
f1 = f1_score(y_test, y_pred_rf, average="weighted")

print('Precision: %0.3f' % precision)
print('Recall: %0.3f' % recall)
print('F1 score: %0.3f' % f1)

Precision: 0.985
Recall: 0.985
F1 score: 0.985
```

```
In [43]: cm = confusion_matrix(y_test, y_pred_rf)
cm_df = pd.DataFrame(cm,
                      index = ['benign', 'defacement', 'phishing', 'malware'],
                      columns = ['benign', 'defacement', 'phishing', 'malware'])
plt.figure(figsize=(6,4))
sns.heatmap(cm_df, annot=True, fmt=".1f")
plt.title('Confusion Matrix')
plt.ylabel('Actual Values')
plt.xlabel('Predicted Values')
plt.show()
```





## Gaussian Naive Bayes

```
In [44]: from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train, y_train)
y_pred_gnb = gnb.predict(X_test)

print(classification_report(y_test, y_pred_gnb, target_names=['benign', 'defacement', 'phishing', 'malware']))

score = accuracy_score(y_test, y_pred_gnb)
print("accuracy:  %0.3f" % score)
```

	precision	recall	f1-score	support
benign	0.85	0.89	0.87	20323
defacement	0.54	1.00	0.70	4742
phishing	0.69	0.44	0.54	3387
malware	0.55	0.12	0.20	4588
accuracy			0.75	33040
macro avg	0.66	0.61	0.58	33040
weighted avg	0.75	0.75	0.72	33040

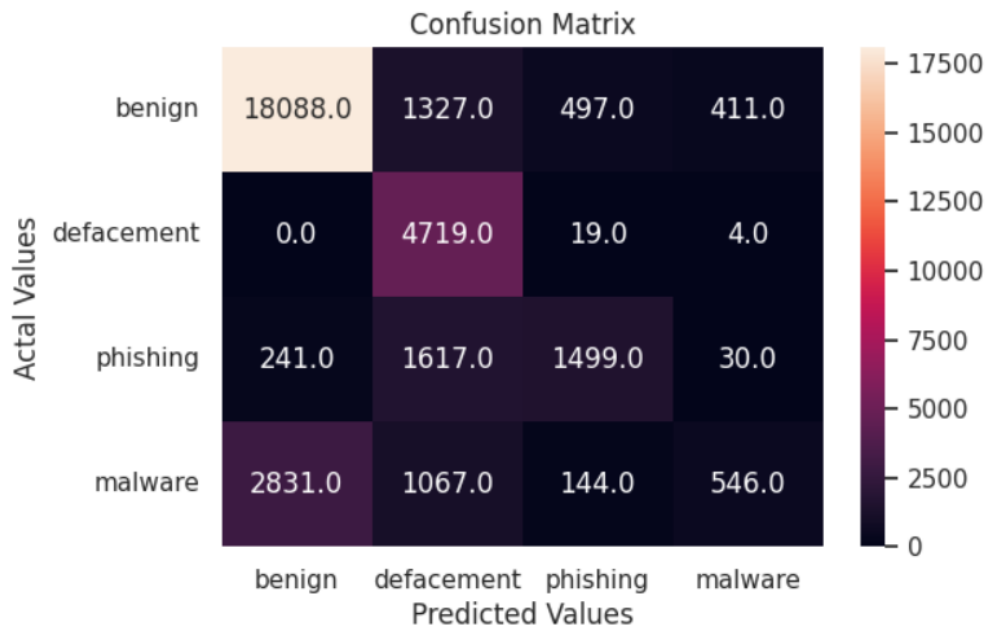
accuracy: 0.752

```
In [58]: # calculate precision, recall, and F1 score
from sklearn.metrics import precision_score, recall_score, f1_score
precision = precision_score(y_test, y_pred_gnb, average="weighted")
recall = recall_score(y_test, y_pred_gnb, average="weighted")
f1 = f1_score(y_test, y_pred_gnb, average="weighted")

print('Precision:  %0.3f' % precision)
print('Recall: %0.3f' % recall)
print('F1 score: %0.3f' % f1)
```

Precision: 0.751  
Recall: 0.752  
F1 score: 0.720

```
In [45]: cm = confusion_matrix(y_test, y_pred_gnb)
cm_df = pd.DataFrame(cm,
                      index = ['benign', 'defacement', 'phishing', 'malware'],
                      columns = ['benign', 'defacement', 'phishing', 'malware'])
plt.figure(figsize=(6,4))
sns.heatmap(cm_df, annot=True, fmt=".1f")
plt.title('Confusion Matrix')
plt.ylabel('Actual Values')
plt.xlabel('Predicted Values')
plt.show()
```



## K Nearest Neighbors

```
In [46]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)

y_pred_knn = knn.predict(X_test)

print(classification_report(y_test, y_pred_knn, target_names=['benign', 'defacement', 'phishing', 'malware']))

score = accuracy_score(y_test, y_pred_knn)
print("accuracy:  %.3f" % score)
```

	precision	recall	f1-score	support
benign	0.99	0.99	0.99	20323
defacement	0.90	0.95	0.92	4742
phishing	0.93	0.87	0.90	3387
malware	0.92	0.89	0.91	4588
accuracy			0.96	33040
macro avg	0.94	0.93	0.93	33040
weighted avg	0.96	0.96	0.96	33040

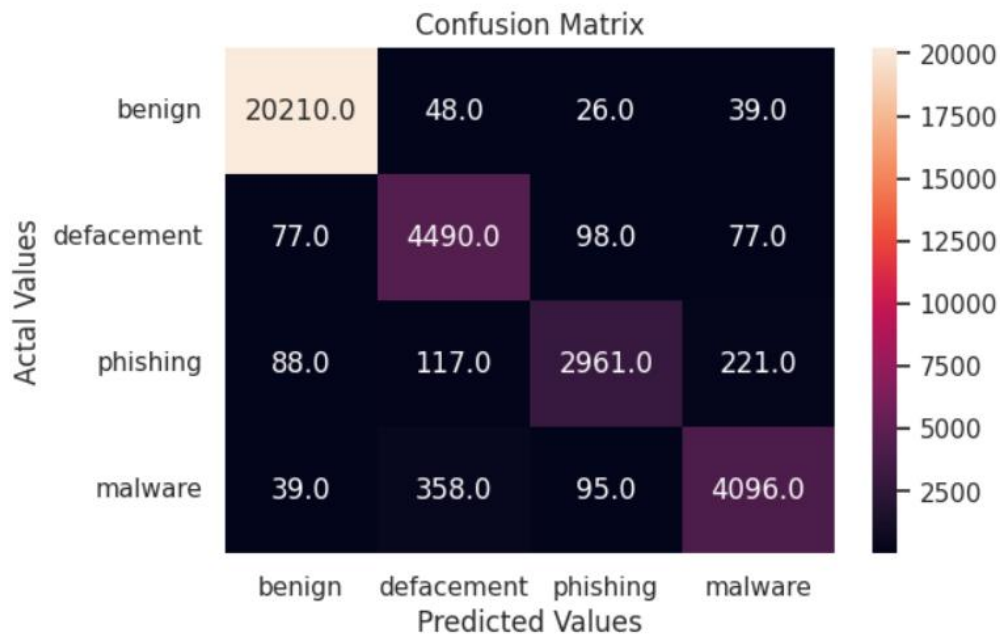
accuracy: 0.961

```
In [59]: # calculate precision, recall, and F1 score
from sklearn.metrics import precision_score, recall_score, f1_score
precision = precision_score(y_test, y_pred_knn, average="weighted")
recall = recall_score(y_test, y_pred_knn, average="weighted")
f1 = f1_score(y_test, y_pred_knn, average="weighted")

print('Precision: %0.3f' % precision)
print('Recall: %0.3f' % recall)
print('F1 score: %0.3f' % f1)

Precision: 0.961
Recall: 0.961
F1 score: 0.961
```

```
In [47]: cm = confusion_matrix(y_test, y_pred_knn)
cm_df = pd.DataFrame(cm,
                      index = ['benign', 'defacement', 'phishing', 'malware'],
                      columns = ['benign', 'defacement', 'phishing', 'malware'])
plt.figure(figsize=(6,4))
sns.heatmap(cm_df, annot=True, fmt=".1f")
plt.title('Confusion Matrix')
plt.ylabel('Actual Values')
plt.xlabel('Predicted Values')
plt.show()
```



## Logistic Regression ¶

```
In [48]: from sklearn.linear_model import LogisticRegression

logreg = LogisticRegression(multi_class='multinomial', solver='lbfgs', max_iter=7000)
logreg.fit(X_train, y_train)
y_pred_lr = logreg.predict(X_test)

print(classification_report(y_test, y_pred_lr, target_names=['benign', 'defacement', 'phishing', 'malware']))

score = accuracy_score(y_test, y_pred_lr)
print("accuracy: %0.3f" % score)
```

	precision	recall	f1-score	support
benign	0.97	0.97	0.97	20323
defacement	0.78	0.89	0.83	4742
phishing	0.76	0.69	0.73	3387
malware	0.87	0.79	0.83	4588
accuracy			0.91	33040
macro avg	0.85	0.84	0.84	33040
weighted avg	0.91	0.91	0.91	33040

accuracy: 0.907

```
In [62]: # calculate precision, recall, and F1 score
from sklearn.metrics import precision_score, recall_score, f1_score
precision = precision_score(y_test, y_pred_lr, average="weighted")
recall = recall_score(y_test, y_pred_lr, average="weighted")
f1 = f1_score(y_test, y_pred_lr, average="weighted")


print('Precision: %0.3f' % precision)
print('Recall: %0.3f' % recall)
print('F1 score: %0.3f' % f1)
```

Precision: 0.907

Recall: 0.907

F1 score: 0.906

```
In [49]: cm = confusion_matrix(y_test, y_pred_lr)
cm_df = pd.DataFrame(cm,
                      index = ['benign', 'defacement', 'phishing', 'malware'],
                      columns = ['benign', 'defacement', 'phishing', 'malware'])
plt.figure(figsize=(6,4))
sns.heatmap(cm_df, annot=True, fmt=".1f")
plt.title('Confusion Matrix')
plt.ylabel('Actual Values')
plt.xlabel('Predicted Values')
plt.show()
```



## Decision Tree

```
In [50]: from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)
y_pred_dt = dt.predict(X_test)

print(classification_report(y_test, y_pred_dt, target_names=['benign', 'defacement', 'phishing', 'malware']))

score = accuracy_score(y_test, y_pred_dt)
print("accuracy:  %0.3f" % score)
```

	precision	recall	f1-score	support
benign	1.00	1.00	1.00	20323
defacement	0.96	0.97	0.96	4742
phishing	0.95	0.92	0.94	3387
malware	0.93	0.95	0.94	4588
accuracy			0.98	33040
macro avg	0.96	0.96	0.96	33040
weighted avg	0.98	0.98	0.98	33040

accuracy: 0.978

```
In [63]: # calculate precision, recall, and F1 score
from sklearn.metrics import precision_score, recall_score, f1_score
precision = precision_score(y_test, y_pred_dt, average="weighted")
recall = recall_score(y_test, y_pred_dt, average="weighted")
f1 = f1_score(y_test, y_pred_dt, average="weighted")


print('Precision: %0.3f' % precision)
print('Recall: %0.3f' % recall)
print('F1 score: %0.3f' % f1)
```

Precision: 0.978

Recall: 0.978

F1 score: 0.978

```
In [51]: cm = confusion_matrix(y_test, y_pred_dt)
cm_df = pd.DataFrame(cm,
                      index = ['benign', 'defacement', 'phishing', 'malware'],
                      columns = ['benign', 'defacement', 'phishing', 'malware'])
plt.figure(figsize=(6,4))
sns.heatmap(cm_df, annot=True, fmt=".1f")
plt.title('Confusion Matrix')
plt.ylabel('Actual Values')
plt.xlabel('Predicted Values')
plt.show()
```





## Support Vector Machines

```
In [52]: from sklearn.svm import SVC
from sklearn.multiclass import OneVsOneClassifier, OneVsRestClassifier
ovo_model = OneVsOneClassifier(SVC())
ovo_model.fit(X_train, y_train)

y_pred_ovm = ovo_model.predict(X_test)

print(classification_report(y_test, y_pred_ovm, target_names=['benign', 'defacement', 'phishing', 'malware']))

score = accuracy_score(y_test, y_pred_ovm)
print("accuracy: %0.3f" % score)
```

	precision	recall	f1-score	support
benign	0.98	0.99	0.99	20323
defacement	0.77	0.93	0.84	4742
phishing	0.83	0.66	0.73	3387
malware	0.90	0.79	0.84	4588
accuracy			0.92	33040
macro avg	0.87	0.84	0.85	33040
weighted avg	0.92	0.92	0.92	33040


accuracy: 0.920

```
In [61]: # calculate precision, recall, and F1 score
from sklearn.metrics import precision_score, recall_score, f1_score
precision = precision_score(y_test, y_pred_ovm, average="weighted")
recall = recall_score(y_test, y_pred_ovm, average="weighted")
f1 = f1_score(y_test, y_pred_ovm, average="weighted")

print('Precision: %0.3f' % precision)
print('Recall: %0.3f' % recall)
print('F1 score: %0.3f' % f1)
```

Precision: 0.922  
Recall: 0.920  
F1 score: 0.918

```
In [53]: cm = confusion_matrix(y_test, y_pred_ovm)
cm_df = pd.DataFrame(cm,
                      index = ['benign', 'defacement', 'phishing', 'malware'],
                      columns = ['benign', 'defacement', 'phishing', 'malware'])
plt.figure(figsize=(6,4))
sns.heatmap(cm_df, annot=True, fmt=".1f")
plt.title('Confusion Matrix')
plt.ylabel('Actual Values')
plt.xlabel('Predicted Values')
plt.show()
```



## Classifier Accuracy Comparison

```
In [64]: # Define the classifiers and their corresponding accuracy scores
classifiers = ['RF', 'GNB', 'KNN', 'LR', 'DT', 'SVM']
accuracy = [0.985, 0.752, 0.961, 0.907, 0.978, 0.92]
precision = [0.985, 0.751, 0.961, 0.907, 0.978, 0.922]
recall = [0.985, 0.752, 0.961, 0.907, 0.978, 0.92]
f1_score = [0.985, 0.720, 0.961, 0.906, 0.978, 0.918]

# Set the width of the bars
barwidth = 0.2

# Set the position of the bars on the x-axis
r1 = range(len(classifiers))
r2 = [x + barwidth for x in r1]
r3 = [x + barwidth for x in r2]
r4 = [x + barwidth for x in r3]

# Create the bar plot
plt.bar(r1, accuracy, color='r', width=barwidth, edgecolor='white', label='Accuracy')
plt.bar(r2, precision, color='g', width=barwidth, edgecolor='white', label='Precision')
plt.bar(r3, recall, color='b', width=barwidth, edgecolor='white', label='Recall')
plt.bar(r4, f1_score, color='c', width=barwidth, edgecolor='white', label='F1 score')

# Add xticks on the middle of the group bars
plt.xlabel('Classifier', fontweight='bold')
plt.xticks([r + barwidth for r in range(len(classifiers))], classifiers)

# Add a legend and title
plt.legend(loc='upper center', bbox_to_anchor=(0.5, -0.15), ncol=4)
plt.title('Evaluation metrics for different classifiers')

# Show the plot
plt.show()
```

